

ETC5521: Diving Deeply into Data Exploration

Working with a single variable, making transformations, detecting outliers, using robust statistics

Professor Di Cook

Department of Econometrics and Business Statistics

Quantitative variables

Outline

- Features to read from a plot of a single quantitative example
- Variety of ways to display a single variable and the pros and cons
- Methods to compute null samples
- Numerical summaries
- Transformations (to symmetrise)
- Handling missing values
- Plotting a single categorical variable
- Ordering categories

Features of a single quantitative variable

Feature	Example	Description
Asymmetry		The distribution is not symmetrical.
Outliers		Some observations are that are far from the rest.
Multimodality		There are more than one "peak" in the observations.
Gaps		Some continuous interval that are contained within the range but no observations exists.

Feature	Example	Description
Heaping		Some values occur unexpectedly often.
Discretized		Only certain values are found, e.g. due to rounding.
Implausible		Values outside of plausible or likely range.

- Is the small bump a set of electorates that could be targeted by Greens in the next election?

Example: election ^(1/5)

First preference votes for the Greens from the 2019 Federal election

► Code

:

Patterns

Analysis

- What is the pattern?
- What are the implications for an analysis?
- Asymmetric: Skewed right
- One outlier
- Multimodal (?)
- Greens won one seat in this election
- To have a chance of winning an electorate a party needs to be one of the top first preferences %. Greens are not close on any other electorate.

Example: wordle ^(2/5)

Distribution of my daily wordle attempts

► Code

:

Patterns

Analysis

- What is the pattern?
- What are the implications for an analysis?
- symmetric (?)
- barrier (?)
- Distribution of rows needed to solve puzzle, likely cannot be a binomial?

Example: olive oil (3/5)

% composition of eicosenoic acid in Italian olive oil samples

► Code

-
- Patterns
- Analysis

- What is the pattern?
- What are the implications for an analysis?
- gap
- Another variable might explain two groups

Example: air quality (4/5)

Concentration of SO2 at sensors across Melbourne throughout 2020.

► Code


-
- Patterns
- Analysis

- What is the pattern?
- What are the implications for an analysis?
- discrete: mix of two distributions
- Resolution for analysis may need to be reduced to the lowest

Example: air travel (5/5)


Air traffic in and out of New York City airports

► Code





-
- Patterns
- Analysis

- What is the pattern?
- What are the implications for an analysis?
- outliers
- symmetric (?)
- implausible (?)
- 50% of observations are within a few minutes of scheduled departure time
- Check the values of the extremes
- Maybe bin up the data, or remove values close to 0, so that the distribution of delayed flights can be examined




► Code

Plot type can affect detectability



► Code



Generating null samples

Activity: Movies

Your turn

1. Grab a copy of the movies data

```
1 data(movies, package = "ggplot2movies")
```

2. Make a plot of the **length** of the movies

3. Describe the structure

4. Play with the type of plot, and the scale, and describe some surprising aspects of movie length

07:35

Simulation

- Normal/Gaussian: `rnorm(n, mean = 0, sd = 1)`
- Uniform: `runif(n, min = 0, max = 1)`
- Exponential: `rexp(n, rate = 1)`
- Log-normal: `rlnorm(n, meanlog = 0, sdlog = 1)`
- Gamma: `rgamma(n, shape, rate = 1)`
- Beta: `rbeta(n, shape1, shape2)`
- Chi-square: `rchisq(n, df)`
- Student's t: `rt(n, df)`
- Weibull: `rweibull(n, shape, scale = 1)`
- Cauchy: `rcauchy(n, location = 0, scale = 1)`
- Binomial: `rbinom(n, size, prob)`
- Negative binomial: `rnbinom(n, size, prob)`
- Geometric: `rgeom(n, prob)`

To generate null samples that might match your data, try `fitdistrplus`:

- `fit <- fitdist(your_data, "norm")`
- `params <- fit$estimate`
- `simulated <- rnorm(n, mean = params["mean"], sd = params["sd"])`

Or compute summary statistics:

Normal

```
mu <- mean(your_data)
sigma <- sd(your_data)
```

Exponential

```
rate <- 1 / mean(your_data)
```

Example: wordle (2/2)

[Binomial](#)

[Learn](#)

► Code

The distribution of actual wordle results is different from samples from a binomial

- counts for 1-2 rows is too low
- peak at 4 is too high

Example: election (1/2)

[Exponential](#)

[Beta](#)

[Learn](#)

► Code

► Code

- The bump is *may not be unusual* for a sample from a standard distribution.
- The single outlier is not typically seen in simulated samples.

Numerical measures of a single quantitative variables

Central tendency and dispersion

- A measure of **central tendency**, e.g. mean, median and mode
- A measure of **dispersion** (also called variability or spread), e.g. variance, standard deviation and interquartile range
- There are other measures, e.g. **skewness** and **kurtosis** that measures “tailedness”, but these are not as common as the measures of first two
- The mean is also the *first moment* and variance, skewness and kurtosis are *second, third, and fourth central moments*

Robust measure of central tendency

- **Mean** is a non-robust measure of location.
- **Median** is the 50% quantile of the observations
- **Trimmed mean** is the sample mean after discarding observations at the tails.
- **Winsorized mean** is the sample mean after replacing observations at the tails with the minimum or maximum of the observations that remain.

Both trimmed and Winsorized mean trimmed 20% of the tails.

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion




25/100

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

26/100




24/08/2026, 09:54

ETC5521: Diving Deeply into Data Exploration



24/08/2026, 09:54

ETC5521: Diving Deeply into Data Exploration



file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

27/100

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

28/100

Plot	Mean	Median	Trimmed Mean	Winsorized Mean
1	0.109	0.114	0.120	0.103
2	0.054	-0.045	-0.016	-0.029
3	1.177	0.729	0.820	0.888
4	0.533	0.541	0.543	0.542
5	0.468	0.329	0.355	0.390
6	5.626	6.656	5.918	5.688



Robust measure of dispersion

- Standard deviation or its square, variance, is a popular choice of measure of dispersion but is not robust to outliers
- Standard deviation for sample $x_1, \dots, x_n$ is


$$\sqrt{\sum_{i=1}^n \frac{(x_i - \bar{x})^2}{n - 1}}$$

- Interquartile range difference between 1st and 3rd quartile, more robust measure of spread
- Median absolute deviance (MAD) is even more robust

$$\text{median}(|x_i - \text{median}(x_i)|)$$



Measure of dispersion					
Plot	SD	IQR	MAD	Skewness	Kurtosis
1	0.90	1.19	0.87	-0.072	3.0
2	0.99	1.41	1.08	0.358	2.2
3	1.33	1.18	0.79	1.944	7.2
4	0.29	0.45	0.34	-0.126	1.8
5	0.47	0.50	0.34	1.691	6.4
6	2.78	5.36	2.98	-0.351	1.7



- Why are the means and the medians different?
- How are the standard deviations and the interquartile ranges similar or different?

Example: election (1/3)



R

% of first preference for the Greens						
Mean	Median	SD	MAD	IQR	Skewness	Kurtosis
9.9	8.5	5.6	3.8	5	2.7	16

- Are there some other numerical statistics we should show?
 - Code

Example: olive oil (2/3)




R

% composition of eicosenoic acid						
Mean	Median	SD	MAD	IQR	Skewness	Kurtosis
16	17	14	21	26	0.34	1.9

- What do the skewness and kurtosis values suggest about the shape of the distribution?

► Code




- Why are the mean and the median so different?
- How do the standard deviation and the interquartile range compare?

Example: air quality (3/3)



R

SO2 concentration (ppm)						
Mean	Median	SD	MAD	IQR	Skewness	Kurtosis
0.00063	0.00041	0.00074	0.00037	0.00051	2.5	9.7



- Why are the mean and the median so different?
- How do the standard deviation and the interquartile range compare?

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion


41/100

Inference for robust statistics

We have seen the [re-sampling methods](#) simulation and permutation used for generating null plots in a lineup. Re-sampling methods can be used with numeric statistics also.

Simulation from distribution, can be used to to check for outliers.

► Code



file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

43/100

- What do the skewness and kurtosis values suggest about the shape of the distribution?

► Code

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

42/100

We can also compute how many simulated values are more than the observed which gives a [simulation \$p\$ -value](#): 0.

For sample [means](#), [conventional tests](#) provide a means for assessing what might be observed if different samples were taken.

[Bootstrapping](#) the current sample, can be used for [robust statistics](#). If we have a sample of values:

```
[1] 2 2 3 6 7 7 8 8
```

to bootstrap sample with replacement:

► Code

```
[1] 2 2 3 3 7 7 7 7
```

► Code

```
[1] 2 3 6 6 6 6 8 8
```

Here's an example of bootstrapping to get a confidence interval for a median.

```
[1] "Median: 6.34"
```

```
[1] "95% CI: ( 4.99 , 9.16 )"
```

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

44/100

 Activity: Air quality

 Your turn

- 1. Look at the **pm25** variable in the **melb_air** data (fine particulate matter, $\mu\text{g}/\text{m}^3$)
- 2. Compute both robust (**median**, **MAD**, **IQR**) and non-robust (**mean**, **SD**) summary statistics – how much do they disagree, and why?
- 3. Bootstrap a 95% confidence interval for the median
- 4. How might these quantities be used to determine reasonable levels of PM25, and indicate extreme values?


05:15

A comment about boxplots (1/2)



data R

► Code



Where are these electorates?

The width of the boxplot is proportional to the number of electoral districts in the corresponding state (which is roughly proportional to the population).


► Code

A comment about boxplots (2/2)



data R

Code



Now what do you notice from this graph that you didn’t notice before?

Transformations

- Only two electoral districts in NT.
- And only 3 and 5 electoral districts in ACT and TAS, respectively!
- Boxplots requires 5 points!
- We should have summarised the number of electoral districts for each state with numerical statistics as a first step.
- Also the outlier (yes, safe to call this an outlier!) and the cluster in the Victoria electorates.

Code

Box-Cox

The Box-Cox transformation is a family of power transformations, indexed by a parameter λ , used to make a variable more symmetric (and closer to normally distributed):

$$y^{(\lambda)} = \begin{cases} \frac{y^\lambda - 1}{\lambda}, & \lambda \neq 0 \\ \log(y), & \lambda = 0 \end{cases}$$

- Only defined for $y > 0$
- λ can be chosen to maximise the log-likelihood of the transformed values being normal, e.g. MASS: `boxcox()`
- The transformation is a guide: round λ to a nearby, interpretable value (e.g. 0 for log, 0.5 for $\sqrt{}$)

The ladder of transformations

Tukey’s **ladder of power transformations** arranges the Box-Cox family by λ . Moving *down* the ladder pulls in long right tails; moving *up* the ladder pulls in long left tails.

λ	Transformation	Typical use
2	y^2	strongly left-skewed
1	y (no change)	roughly symmetric
0.5	$\sqrt{y}$	mild right-skew, count data
0	$\log(y)$	moderate to strong right-skew
-0.5	$1/\sqrt{y}$	strong right-skew
-1	$1/y$	very strong right-skew, e.g. rates

- Start close to $\lambda = 1$ (no transformation) and move up or down the ladder, checking symmetry with a plot after each step
- The further λ moves from 1, the stronger the transformation

Example: air quality (PM2.5)




data **R**

Original **pm25** ($\mu\text{g}/\text{m}^3$) – 34 zero values are **shifted by** $+0.1$ before transforming, since Box-Cox requires $y > 0$:

► Code


Box-Cox log-likelihood profile:

► Code



Optimal $\lambda \approx 0.2$
Transformed variable, using the optimal λ directly (not rounded to a ladder rung):

► Code




- Skewness before transforming: 11.04
- Skewness after transforming: -0.17
- $\lambda \approx 0.2$ sits between $\sqrt{y}$ ($\lambda = 0.5$) and $\log(y)$ ($\lambda = 0$) on the ladder. Rounding to $\log(y)$ actually *overcorrects* here (skewness becomes -1.33, i.e. left-skewed) – a reminder that the optimal λ is not always well-approximated by a “nice” rung when there are extreme outliers (here, from bushfire smoke).

► Code

► Code

Non box-cox example: Olive oil ^(1/2)

► Code




24/08/2026, 09:54

ETC5521: Diving Deeply into Data Exploration


Two clusters

Check if there is another variable that explains the two groups


► Code



Non box-cox example: Olive oil (2/2)



Mixture of discreteness and normal shape of continuous values. *Why might this happen?*
Check if there is a difference in the strata (here 1 thru 9), implying measurement policy differences.



 Activity: Air quality

 Your turn

When and why do you want to make a transformation of a variable?

02:01

Suburb	Rooms	Type	Price (\$)	Date
Airport West	2	Unit	540,000	2017-04-01
Airport West	3	Home	715,000	2017-04-01
Albanvale	6	Home	NA	2017-04-01
Albert Park	3	Home	1,925,000	2017-04-01
Albion	3	Unit	515,000	2017-04-01
Albion	4	Home	717,000	2017-04-01
Alphington	2	Home	1,675,000	2017-04-01
Alphington	4	Home	2,008,000	2017-04-01

Melbourne Housing Prices ^(1/5)

► Code

Suburb	Rooms	Type	Price (\$)	Date
Abbotsford	3	Home	1,490,000	2017-04-01
Abbotsford	3	Home	1,220,000	2017-04-01
Abbotsford	3	Home	1,420,000	2017-04-01
Aberfeldie	3	Home	1,515,000	2017-04-01
Airport West	2	Home	670,000	2017-04-01
Airport West	2	Townhouse	530,000	2017-04-01

Suburb	Rooms	Type	Price (\$)	Date
Altona	2	Home	860,000	2017-04-01
Altona Meadows	4	Home	NA	2017-04-01
Altona North	3	Home	720,000	2017-04-01
Armadaale	2	Unit	836,000	2017-04-01
Armadaale	2	Home	2,110,000	2017-04-01
Armadaale	3	Home	1,386,000	2017-04-01

- This data was scraped each week from domain.com.au from 2016-01-28 to 2018-10-13
- In total there are **63,023** observations


- All variables shown (there are more variables not shown here), except price, have complete records
- There are **48,433** property prices across Melbourne (roughly 23% missing)

Data source: Tony Plo (2018) [Melbourne Housing Market](#)

How would you explore this data first?


Yes, with an overview plot.

Melbourne Housing Prices (2/5)



There are only missings on price, but a lot of them. Can we find some other collected variable which might help to explain the missing price?
Is missingness more likely for expensive houses?

Use a substitute like **Rooms** to examine the possible relationship. Show **Rooms** vs missing on **Price** and use a **lineup** before looking at the data.
Rooms is discrete, so a bar chart, possibly faceted or side-by-side bars, would be a suitable plot.



- To impute missings other variables will need to be used.

Note: Houses with more than 8 rooms removed. Why?


► Code

► Code

Data Summary					
Name	df2				
Number of rows	63023				
Number of columns	13				
Column type frequency:					
character	8				
Date	1				
numeric	4				
Group variables	None				
Variable type: character					
skim_variable	n_missing	complete_rate	min	max	empty
1 Suburb	0	1	3	18	0
2 Address	0	1	7	27	0
3 Type	0	1	1	1	0
4 Method	0	1	1	2	0
5 SellerG	0	1	1	27	0
6 Postcode	0	1	4	4	0
7 Regionname	0	1	16	26	0

► Code

► Code



Is there a suspicious plot?

Break the association between Rooms and Missing/Not on Price, because the null hypothesis is that there is **no difference in missing status for price based on the size of the house**. Why?

► Code

Check the support of your data

If you have **too few measurements** in any region (extreme), summaries for these regions will be unreliable.


- For quantitative variables, it may be necessary to **remove extremes**.
- If the variable is categorical it might be best to **combine levels**.
- It is important to **script** so decisions can be reversed or **rare events** are not ignored.

We removed houses with 8 or more rooms. What other way might we have handled these houses?

Melbourne Housing Prices (3/5)



data R



What can we say from this plot?

- The housing prices are right-skewed
- There appears to be a lot of outlying housing prices (how can we tell?)

Note: We determined that it is likely that more higher price houses have not disclosed the sale price. The distribution of price will need to be checked again **after imputation**.

► Code

► Code


► Code

► Code

Melbourne Housing Prices (4/5)



data R



- The x-axis has been $\log_{10}$ -transformed in this plot
- The plot appears more symmetrical now
- What is a useful measure of central tendency here?

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

81/100

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

82/100

Melbourne Housing Prices (5/5)

Central tendency R

With no transformation:

Mean	Median	Trimmed Mean	Winsorised Mean
\$997,898	\$830,000	\$871,375	\$903,823

With log transformation (and back-transformed to original scale):

Mean	Median	Trimmed Mean	Winsorised Mean
\$874,166	\$830,000	\$847,973	\$859,325

► Code

► Code

Categorical variables

There are two types of categorical variables

Nominal where there is no intrinsic ordering to the categories
E.g. blue, grey, black, white.

Ordinal where there is a clear order to the categories.
E.g. Strongly disagree, disagree, neutral, agree, strongly agree.

Levels: 1 3 10

► Code

[1] 1 3 10
Levels: 1 10 3

► Code

[1] 1 3 10
Levels: 1 3 10

Categorical variables in R

- In R, categorical variables may be encoded as **factors**.

► Code

```
[1] 2 2 1 1 3 3 3 1  
Levels: 1 2 3
```

- You can easily change the labels of the variables:

► Code

```
[1] II II I I III III III I  
Levels: I II III
```

- Order of the factors are determined by the input:

► Code

```
[1] 1 3 10
```

Order nominal variables meaningfully

Coding tip: use below functions to easily change the order of factor levels

```
1 stats::reorder(factor, value, mean)  
2 forcats::fct_reorder(factor, value, median)  
3 forcats::fct_reorder2(factor, value1, value2, func)
```

Numerical summaries

counts, proportions, percentages and odds

Tuberculosis counts in Australia

► Code

```
# A tibble: 22 × 7
  country iso3 year count      p    pct odds
  <chr>   <chr> <dbl> <dbl> <dbl> <dbl> <dbl>
1 Australia AUS 2000   982 0.0522 5.22 1
2 Australia AUS 2001   953 0.0507 5.07 0.970
3 Australia AUS 2002  1008 0.0536 5.36 1.03
4 Australia AUS 2003   926 0.0493 4.93 0.943
5 Australia AUS 2004  1036 0.0551 5.51 1.05
6 Australia AUS 2005  1030 0.0548 5.48 1.05
7 Australia AUS 2006  1127 0.0600 6.00 1.15
8 Australia AUS 2007  1081 0.0575 5.75 1.10
9 Australia AUS 2008  1182 0.0629 6.29 1.20
10 Australia AUS 2009  1176 0.0626 6.26 1.20
11 Australia AUS 2010  1146 0.0610 6.10 1.17
12 Australia AUS 2011  1202 0.0640 6.40 1.22
13 Australia AUS 2012  1259 0.0670 6.70 1.28
14 Australia AUS 2013   512 0.0272 2.72 0.521
15 Australia AUS 2014   474 0.0252 2.52 0.483
16 Australia AUS 2015   438 0.0233 2.33 0.446
17 Australia AUS 2016   481 0.0256 2.56 0.490
18 Australia AUS 2017   524 0.0279 2.79 0.534
19 Australia AUS 2018   502 0.0267 2.67 0.511
```

For qualitative data, compute

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

89/100

- count/frequency,
- proportion/percentage
- and sometimes, an **odds ratio**. Here we have used ratio relative to the count in year 2000.


Note: For exploration, no rounding of digits was done, but to report you would need to make the numbers pretty.

Ways to plot a single categorical variable

► Code

file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

91/100




file:///Users/cookd/teaching/ETC5521/ddde-private/docs/week5/slides.html#/robust-measure-of-dispersion

92/100



Number of attempts to solve Wordle:

- **Bar chart:** height encodes count directly – easiest to compare categories accurately
 - **Pie chart:** angle (and area) encodes count – comparing categories by angle is harder than by length
 - **Rose plot** (nightingale/coxcomb): like a bar chart wrapped into a circle – radius encodes count, but the outer wedges look larger than their true count because area grows with the square of the radius
 - **Spine plot:** a single bar split into proportions – good for comparing proportions to the whole, or as a building block for two categorical variables (stacking within categories of a second variable)
- Generally, **bar charts** (or spine plots) are the easiest to read accurately: humans compare lengths/positions along a common scale far better than angles or areas.



ETCS521: Diving Deeply into Data Exploration

24/08/2026, 09:54


Activity: Melbourne housing

💡 Your turn

- Make the spine plot of **Rooms** and fill the bars based on missingness on Price.
- If you can do this using a lineup, even better.

The code on slide 35 might help.

05:52



Categorical variable: Simulate from multinomial.

► Code

```
# A tibble: 7 × 2
  age count
<chr> <dbl>
1 1524    27
2 2534    40
3 3544    15
4 4554    11
5 5564     9
6 65     15
7 u      12
```

► Code

```
[1] 5 3 1 1 2 1 2 1 1 1 1
```

► Code

```
# A tibble: 7 × 2
  age count
<chr> <dbl>
1 1524    35
2 2534    50
3 3544    16
```


4	4554	11
5	5564	10
6	65	15
7	u	12

`imputeMulti` library can automate for multiple variables.

- For a categorical variable, compute [counts](#), [proportions/percentages and odds](#), and order categories meaningfully rather than relying on the default ordering

Key points

- [Make many plots](#) of a single variable (histogram, density, boxplot, quasirandom) – the type of plot, and choices like bin width or bandwidth, change what patterns are visible
- Use [simulation and permutation](#) to generate null samples and lineups, to judge whether a striking feature (outlier, cluster, shape) is more than what’s expected by chance
- Compute both non-robust ([mean](#), [SD](#)) and robust ([median](#), [MAD](#), [IQR](#), [trimmed/winsorised mean](#)) statistics – skew and outliers can pull the non-robust ones a long way from the bulk of the data
- Boxplots are a compact robust summary, but can hide multimodality – check against the raw data
- [Transformations](#) (e.g. Box-Cox, the ladder of power transformations) can re-focus attention and symmetrise skewed data, but round the “optimal” λ with care, especially when extreme outliers are present
- Check for [missing values](#) early: visualise the pattern of missingness (e.g. `vis_miss()`, a lineup) before deciding how, or whether, to impute – and script decisions so they can be reversed

Resources

- Unwin (2015) Graphical Data Analysis with R
- Venables, W. N. & Ripley, B. D. (2002) Modern Applied Statistics with S. Fourth Edition. Springer, New York. ISBN 0-387-95457-0
- Josse et al (2022) [R-miss-tastic](#)
- Wilke [Fundamentals of Data Visualization](#), chapters 6-11.



ETC5521 Lecture 5 | [ddde.numbat.space](#)

Speaker notes